

Conjugate models 02

Josep Fortiana

Facultat de Matemàtiques i Informàtica, UB

2021-03-15

1 The Poisson-Gamma model

Likelihood

The likelihood is:

$$(y|\lambda) \sim \text{Poisson}(\lambda),$$

with pmf:

$$f(y|\lambda) = e^{-\lambda} \frac{\lambda^y}{y!}, \quad y = 0, 1, 2, \dots, \quad \lambda > 0,$$

Conjugate prior pdf

λ 's prior is:

$$\text{Gamma}(\alpha, \beta), \quad \alpha, \beta > 0,$$

with pdf:

$$h(\lambda) = \frac{\beta^\alpha}{\Gamma(\alpha)} \lambda^{\alpha-1} e^{-\beta\lambda}, \quad \alpha, \beta, \lambda > 0.$$

Joint pdf

The joint “density” of (y, λ) is:

$$f(y, \lambda) = f(y|\lambda) \cdot h(\lambda) = \frac{\beta^\alpha}{\Gamma(\alpha)} \cdot \frac{\lambda^{\alpha+y-1}}{y!} \exp\{-(\beta+1)\lambda\},$$

for $\alpha, \beta, \lambda > 0$.

Marginal pmf of y (Prior predictive pmf)

To integrate with respect to λ we split $f(y, \lambda)$:

$$\begin{aligned} & \frac{\beta^\alpha}{\Gamma(\alpha)} \cdot \frac{1}{y!} \cdot \frac{\Gamma(\alpha+y)}{(\beta+1)^{(\alpha+y)}} \\ & \times \frac{(\beta+1)^{(\alpha+y)}}{\Gamma(\alpha+y)} \cdot \lambda^{\alpha+y-1} \exp\{-(\beta+1)\lambda\}, \end{aligned}$$

The second line is a $\text{Gamma}(\alpha+y, \beta+1)$ pdf, which integrates to 1.

Marginal pmf of y

After integrating out λ we get the marginal pmf of y :

$$f(y) = \frac{\Gamma(\alpha + y)}{\Gamma(\alpha) y!} \frac{\beta^\alpha}{(\beta + 1)^{\alpha+y}}, \quad y = 0, 1, \dots, \quad \alpha, \beta > 0$$

Noting that:

$$\frac{\beta^\alpha}{(\beta + 1)^{\alpha+y}} = \left(\frac{1}{\beta + 1} \right)^y \cdot \left(\frac{\beta}{\beta + 1} \right)^\alpha,$$

Marginal pmf of y

we see $f(y)$ is a $\text{NegBin}(r, p)$ pmf, with parameters:

$$r = \alpha, \quad \text{and} \quad p = \frac{\beta}{\beta + 1}.$$

(see Section [A](#) in the Appendix).

This is the *prior predictive* distribution for y .

Posterior pdf

With Bayes' formula, dividing $f(y, \lambda)$ by the marginal pmf we obtain the posterior pdf of λ , given y (first line above):

$$h(\lambda | y) = \frac{(\beta + 1)^{\alpha+y}}{\Gamma(\alpha + y)} \lambda^{\alpha+y-1} e^{-(\beta+1)\lambda}, \quad \lambda > 0,$$

which is a $\text{Gamma}(\alpha + y, \beta + 1)$ pdf.

Case of an n -sample: posterior pdf

For y_1, \dots, y_n i.i.d. $\sim \text{Poisson}(\lambda)$, the sum:

$$y = \sum_{i=1}^n y_i \sim \text{Poisson}(n \lambda)$$

(Additive property of Poisson r.v.)

Thus, for a prior $\lambda \sim \text{Gamma}(\alpha, \beta)$ and n $\text{Poisson}(\lambda)$ data, the posterior is a $\text{Gamma}(\alpha + y, \beta + n)$.

Case of an n -sample: prior predictive pmf

Similarly, for n observed $\text{Poisson}(\lambda)$ data, the prior predictive (marginal) distribution of the total count number $y = \sum_{i=1}^n y_i$ is a $\text{NegBin}(r, p)$ r.v., with parameters:

$$r = \alpha, \quad \text{and} \quad p = \frac{\beta}{\beta + n}.$$

An n -sample with different exposures

In many applications we find data of the form:

$$y_i \sim \text{Poisson}(\lambda_i), \quad \text{where } \lambda_i = x_i \cdot \theta, \quad 1 \leq i \leq n.$$

The values x_i are known positive values of an explanatory variable, usually called *exposure*, and θ is the common *rate* parameter.

Posterior pdf for an n -sample with different exposures

With a $\theta \sim \text{Gamma}(\alpha, \beta)$ prior,

observations $\mathbf{y} = (y_1, \dots, y_n)$, exposures $\mathbf{x} = (x_1, \dots, x_n)$,

the posterior pdf is:

$$\theta | \mathbf{y} \sim \text{Gamma}(\alpha + y, \beta + x),$$

where $y = \sum_{i=1}^n y_i$ and $x = \sum_{i=1}^n x_i$.

2 The Dirichlet - Multinomial model

The Dirichlet - Multinomial model

Generalizes the Beta-Binomial conjugate pair.

The Dirichlet distribution is the multivariate Beta.

The Multinomial distribution is the multivariate Binomial.

Multivariate Bernoulli distribution

A partition $\Omega = A_1 \sqcup \dots \sqcup A_m$, where the A_j are pairwise exclusive events whose union is the total space, and:

$$\theta = (\theta_1, \dots, \theta_m), \quad \theta_j = P(A_j), \quad 1 \leq j \leq m.$$

Each indicator

$$\mathbb{1}_{A_j} \sim \text{Bernoulli}(\theta_j) \quad 1 \leq j \leq m.$$

Multivariate Bernoulli distribution

The m -dimensional vector of indicators:

$$(\mathbb{1}_{A_1}, \dots, \mathbb{1}_{A_m}),$$

follows an m -dimensional *multivariate Bernoulli* distribution with vector of probabilities:

$$\theta = (\theta_1, \dots, \theta_m),$$

Multivariate Bernoulli distribution

The sum of the m probabilities is 1.

The sum of the m indicators is 1 (**Cannot be independent!!**).

Each j -th marginal, $\mathbb{1}_{A_j} \sim \text{Bernoulli}(\theta_j)$.

$$E(\mathbb{1}_{A_j}) = \theta_j,$$

$$\text{var}(\mathbb{1}_{A_j}) = \theta_j (1 - \theta_j),$$

$$\text{cov}(\mathbb{1}_{A_j}, \mathbb{1}_{A_k}) = -\theta_j \theta_k, \quad j \neq k.$$

The multinomial distribution

The m -dim. multinomial distribution of size n and probs. $\theta = (\theta_1, \dots, \theta_m)$, $\theta_j \in [0, 1]$, $\sum_{j=1}^m \theta_j = 1$, has joint pmf:

$$\frac{n!}{x_1! \dots x_m!} \theta_1^{x_1} \dots \theta_m^{x_m}.$$

for an m -dimensional vector $x = (x_1, \dots, x_m)$, of integers $x_j \in [0, n]$ such that $\sum_{j=1}^m x_j = n$.

x is the sum of n m -dim. vectors i.i.d. $\sim \text{Bernoulli}(\theta)$.

The Dirichlet distribution

θ 's joint pdf, with parameters $a = (a_1, \dots, a_m)$, $a_i > 0$:

$$h(\theta_1, \dots, \theta_m; a_1, \dots, a_m) = \frac{1}{B(a)} \prod_{i=1}^m \theta_i^{a_i-1}, \quad \text{where}$$

$B(a) = \frac{\prod_{i=1}^m \Gamma(a_i)}{\Gamma(\sum_{i=1}^m a_i)}$ is the multivariate Beta function.

Multinomial likelihood with Dirichlet prior

If: $(x|\theta) \equiv (x_1, \dots, x_m | \theta_1, \dots, \theta_m) \sim \text{Multinom}(n, \theta)$,

with parameter vector:

$$\theta = (\theta_1, \dots, \theta_m), \quad 0 < \theta_i < 1, \quad 1 \leq i \leq m, \quad \sum_{i=1}^m \theta_i = 1,$$

and θ 's joint prior is Dirichlet with parameters

$a = (a_1, \dots, a_m)$, then θ 's posterior is Dirichlet,

with parameters: $a + x = (a_1 + x_1, \dots, a_m + x_m)$.

3 Exponential-Gamma model

Exponential likelihood

Used to model *waiting times* or *insurance claims*.

The pdf of an outcome y , given θ , is:

$$f(y|\theta) = \theta \exp(-\theta \cdot y), \quad y > 0,$$

and $\theta = 1/E(y|\theta)$ is called the *rate parameter*.

The exponential is a special case of a Gamma with $(\alpha, \beta) = (1, \theta)$.

Gamma prior for θ

The conjugate prior for θ is a $\text{Gamma}(\alpha, \beta)$, for $\alpha, \beta > 0$, with pdf:

$$h(\theta) = \frac{\beta^\alpha}{\Gamma(\alpha)} \cdot \theta^{\alpha-1} \cdot \exp\{-\beta \theta\}, \quad \theta > 0.$$

Likelihood for an n -sample

$y = (y_1, \dots, y_n)$ i.i.d. $\sim \text{Exp}(\theta)$.

$$f(y|\theta) = \prod_{i=1}^n (\theta e^{-\theta y_i}) = \theta^n \cdot \exp\{-n\theta \bar{y}\}.$$

Joint (y, θ) pdf

$$\begin{aligned} f(y, \theta) &= f(y|\theta) \cdot h(\theta) \\ &= \frac{\beta^\alpha}{\Gamma(\alpha)} \cdot \theta^{\alpha+n-1} \cdot \exp\{-\theta \cdot (n\bar{y} + \beta)\} \end{aligned}$$

To get the $y = (y_1, \dots, y_n)$ joint marginal (prior predictive pdf) we integrate $f(y, \theta)$ with respect to θ .

(Integral is a gamma function, after a variable change $t = \theta (n \bar{y} + \beta)$)

Prior predictive pdf

$$f(y_1, \dots, y_n) = \frac{\Gamma(\alpha + n)}{\Gamma(\alpha)} \cdot \frac{\beta^\alpha}{(n \bar{y} + \beta)^{\alpha+n}},$$

$$y_1 > 0, \dots, y_n > 0.$$



Caution! This is the joint pdf of the y vector.

It is a function of the sum $n \bar{y}$ but NOT the pdf OF $x = n \bar{y}$.
Obtaining it requires $n - 1$ further integrals.

Prior predictive density of $x = n \bar{y}$ for small n

For $n = 1$,

$$f(x) = \frac{\alpha \beta^\alpha}{(x + \beta)^{\alpha+1}}, \quad x > 0,$$

for $n = 2$,

$$f(x) = (\alpha + 1) \alpha \beta^\alpha \frac{x}{(x + \beta)^{\alpha+2}}, \quad x > 0,$$

$\alpha > 0, \beta > 0$.

Prior predictive density of $x = n \bar{y}$ for small n

For $n = 1$,

$$E(x) = \frac{\beta}{\alpha - 1}, \quad \alpha > 1, \quad \text{var}(x) = \frac{\alpha \beta^2}{(\alpha - 2)(\alpha - 1)^2}, \quad \alpha > 2.$$

For $n = 2$,

$$E(x) = \frac{2\beta}{\alpha - 1}, \quad \alpha > 1, \quad \text{var}(x) = \frac{2(\alpha + 1)\beta^2}{(\alpha - 2)(\alpha - 1)^2}, \quad \alpha > 2.$$

$(\beta > 0)$.

Posterior pdf of θ , given y

$$(\theta | y) \sim \text{Gamma}(\alpha', \beta'),$$

where the updating formula is:

$$\begin{cases} \alpha' &= \alpha + n, \\ \beta' &= \beta + n \bar{y}. \end{cases}$$

4 Mixture priors: the spinning coin

Persi Diaconis

Stanford prof. Persi Diaconis, formerly a professional magician, famously found how many shuffles a deck of cards needs to give a truly random order (**seven**). He's also dabbled in coin games.




BTW: Persi Diaconis on randomness

See his 2013 video talk:

[The Search for Randomness](#)

The spinning coin Fact: if a coin is spun on its edge instead of being flipped, proportion of heads or tails is not around 50% but rather such values as 25% or 75% are obtained.



Persi Diaconis on the spinning coin

According to Diaconis, *“the reasons for the bias are not hard to infer. The shape of the edge will be a strong determining factor – indeed, magicians have coins that are slightly shaved; the eye cannot detect the shaving, but the spun coin always comes up heads”*.

A prior for the spinning coin problem

For n tosses of a spinning coin, the number x of heads up is a binomial $B(n, \theta)$, and θ 's prior will typically be bimodal (pdf with two local maxima). It cannot be a $\text{Beta}(\alpha, \beta)$, which has a single mode at:

$$\frac{\alpha - 1}{\alpha + \beta - 2}$$


[Diaconis, Persi and Donald Ylvisaker \(1985\) *Quantifying prior opinion*.](#)

[In:](#)

[J.M. Bernardo et al \(eds\), *Bayesian Statistics 2*, Elsevier, pp. 133-156.](#)

A possible prior

$$0.50 \text{ Beta}(10, 20) + 0.50 \text{ Beta}(20, 10).$$



Interpretation of a mixture prior

The mixture prior can be thought of as a weighted combination of “beta populations”, the weights γ_i measuring the prior degree of belief that the actual coin was chosen from the i -th population.


To generate random numbers from a mixture:

$$\gamma \cdot h_1 + (1 - \gamma) \cdot h_2, \quad \gamma \in (0, 1),$$

we generate a sequence of realizations of $I \sim \text{Bernoulli}(\gamma)$ and:

- For each entry equal to 1, a realization of h_1 ,
- For each entry equal to 0, a realization of h_2 .

Simulation of a mixture prior



Bayesian modelling with a mixture prior

Assume the prior pdf for θ is:

$$h(\theta) = \gamma \cdot h_1(\theta) + (1 - \gamma) \cdot h_2(\theta),$$

and the likelihood is: $f(x|\theta)$. Then, the joint pdf is:

$$f(x, \theta) = \gamma \cdot f_1(x, \theta) + (1 - \gamma) \cdot f_2(x, \theta).$$

where:

$$f_i(x, \theta) = f(x|\theta) \cdot h_i(\theta), \quad i = 1, 2.$$

Prior predictive pdf from a mixture prior

Integrating out θ , the marginal for x is:

$$f(x) = \gamma \cdot f_1(x) + (1 - \gamma) \cdot f_2(x),$$

where:

$$f_i(x) = \int f(x|\theta) \cdot h_i(\theta) d\theta, \quad i = 1, 2.$$

Computing the posterior pdf from a mixture prior

From Bayes' formula:

$$h(\theta|x) = \frac{f(x, \theta)}{f(x)} = \frac{\gamma \cdot f_1(x, \theta) + (1 - \gamma) \cdot f_2(x, \theta)}{\gamma \cdot f_1(x) + (1 - \gamma) \cdot f_2(x)}.$$

With the obvious notation:

$$h_i(\theta|x) = \frac{f_i(x, \theta)}{f_i(x)} \quad i = 1, 2,$$

Posterior pdf from a mixture prior

the posterior pdf is:

$$h(\theta|x) = \hat{\gamma}(x) \cdot h_1(\theta|x) + (1 - \hat{\gamma}(x)) \cdot h_2(\theta|x),$$

where the *posterior mixture weights* are:

$$\hat{\gamma}(x) = \frac{\gamma \cdot f_1(x)}{\gamma \cdot f_1(x) + (1 - \gamma) \cdot f_2(x)}, \text{ and } 1 - \hat{\gamma}(x).$$

A Appendix: The negative binomial distribution

First definition

A sequence of independent binary 0/1 experiments, with i.i.d. $\sim \text{Bernoulli}(p)$ indicators, $p \in (0, 1)$.

Number X of realizations to obtain a number $r \in \mathbb{N}$ of successes (1's), is a *negative binomial r.v.* with size r and probability p .

Alternative: $Y = X - r$ = number of failures (0's) before obtaining a number r of successes.

Probability mass function

$$P(x) = \binom{x-1}{r-1} \cdot (1-p)^{x-r} \cdot p^r, \quad x = 1, 2, \dots,$$

$$P(y) = \binom{y+r-1}{r-1} \cdot (1-p)^y \cdot p^r, \quad y = 0, 1, 2, \dots,$$

The second one is more usual (e.g., `dnbinom()` in R).

Alternative (and the reason for the name)

$$P(y) = \binom{-r}{y} \cdot p^r \cdot (-q)^y, \quad \text{where } q = 1 - p.$$

Indeed:

$$\begin{aligned} \binom{-r}{y} &= \frac{(-r) \cdot (-r-1) \cdot \dots \cdot (-r-y+1)}{y!} \\ &= (-1)^y \cdot \binom{y+r-1}{r-1}. \end{aligned}$$

General definition

For an integer r ,

$$\binom{y+r-1}{r-1} = \frac{(y+r-1)!}{(r-1)! \cdot y!} = \frac{\Gamma(y+r)}{\Gamma(r) \cdot y!}.$$

The right hand is valid for real $r > 0$. Thus the pmf:

$$P(y) = \frac{\Gamma(y+r)}{\Gamma(r) \cdot y!} \cdot (1-p)^y \cdot p^r, \quad y = 0, 1, 2, \dots,$$

defines the $\text{NegBin}(r, p)$, for $r > 0$ and probability p .

Relation to the geometric distribution

For $r = 1$, the $\text{NegBin}(1, p)$ is the $\text{Geom}(p)$.

For integer r , a r.v. distributed as $\text{NegBin}(r, p)$ can be considered as the sum of r i.i.d. copies of a $\text{Geom}(p)$.

Expectation, variance of a negative binomial

For $Y \sim \text{NegBin}(r, p)$,

$$E(Y) = \mu \equiv r \cdot \frac{1-p}{p}$$

$$\text{var}(Y) = \sigma^2 \equiv r \cdot \frac{1-p}{p^2} = \frac{\mu}{p} = \mu + \frac{\mu^2}{r}.$$
