

NUMERICAL LINEAR ALGEBRA  
MASTER ON FUNDAMENTALS OF DATA SCIENCE  
UNIVERSITAT DE BARCELONA

---

# Project 3

## PageRank Implementations

---

Blai Ras

December 21, 2020

# Contents

|          |   |          |
|----------|---|----------|
| <b>1</b> | <b>Results</b>                                  | <b>3</b> |
| 1.1      | Linear system . . . . .                         | 3        |
| 1.2      | Power method storing matrices . . . . .         | 4        |
| 1.3      | Power method without storing matrices . . . . . | 5        |
| 1.4      | Convergence . . . . .                           | 5        |

# 1 Results

The Page Rank algorithm can be implemented in multiple ways. In class, we are introduced to 2 different types of implementations: as a linear system and with the power method. The second one can be programmed in two different ways: storing the matrices and without.

In this section we will be discussing the different outputs of each method with two examples of sparse matrices: `p2p-Gnutella30` (36682 x 36682) and `p2p-Gnutella31` (62586 x 62586), extracted from [the University of Florida Sparse Matrix Collection](#).

## 1.1 Linear system

Trying to find the PR vector of a big sparse matrix is a very bad idea. We want:

$$\hat{M}m = (1 - m)A + \hat{S}$$

$$\text{where } A = G \cdot D$$

$$\text{where } D \text{ is a diagonal matrix as } = \begin{cases} \frac{1}{n_j} & \text{if } n_j \neq 0 \\ 0 & \text{otherwise} \end{cases}$$

$$\text{and } n_j = \sum_i g_{ij}$$

Isolating  $x$  we have:

$$(Id - (1 - m)GD)x = e\hat{z}^T x$$

This can be done with the `scipy.sparse.linalg.spsolve` method if we take in consideration that  $\sum x_i = 1$  and therefore  $b$  is an array of dimension  $(n \times 1)$  of ones.

The computation of the PR vector using this method lasted for one minute for the first matrix and 6.379 minutes for the second one, using a damping factor of 0.15. In conclusion, too much computational cost.

## 1.2 Power method storing matrices

This time we have again:

$$\hat{M}m = (1 - m)A + \hat{S}$$

But

$$\hat{S} = e\hat{z}^T \text{ where } \hat{z}_j = \begin{cases} \frac{m}{n} & \text{if the column } j \text{ of } A \text{ is non-zero} \\ \frac{1}{n} & \text{otherwise} \end{cases}$$

In this case the operation that takes the most computational time is the creation of the  $z_j$  array. This are my results with a damping factor of 0.15:

| Matrix         | Tolerance | Time (s) |
|----------------|-----------|----------|
| p2p-Gnutella30 | 1e-5      | 20.73    |
| p2p-Gnutella30 | 1e-15     | 21.12    |
| p2p-Gnutella31 | 1e-5      | 61.28    |
| p2p-Gnutella31 | 1e-15     | 61.91    |

We can see that the tolerance does not affect much the computational time. Let's try changing the damping factor to 0.35:

| Matrix         | Tolerance | Time (s) |
|----------------|-----------|----------|
| p2p-Gnutella30 | 1e-5      | 20.73    |
| p2p-Gnutella30 | 1e-15     | 21.12    |

Again, the two times are very similar. I avoid the computation for the even bigger p2p-Gnutella31 matrix.

### 1.3 Power method without storing matrices

Same idea but this time we are working with indexes instead of storing the whole matrix. This means considering:

$$j \text{ such that } n_j \neq 0 \Rightarrow M_{ij} \begin{cases} 0 & \text{if } g_{ij} = 0 \\ (1 - m) & \frac{1}{n_j} \text{ if } g_{ij} = 1 \end{cases}$$

$$g_{ij} = 0 \iff \nexists \text{ link } i \rightarrow j \iff i \not\in L_j$$

My results, with a damping factor of 0.15 are:

| Matrix         | Tolerance | Time (s) |
|----------------|-----------|----------|
| p2p-Gnutella30 | 1e-5      | 4        |
| p2p-Gnutella30 | 1e-15     | 4.08     |
| p2p-Gnutella31 | 1e-5      | 8.62     |
| p2p-Gnutella31 | 1e-15     | 8.61     |

We definitely see a big time improvement without the storage of the matrices. I see a bit of improvement if I lower the tolerance. Let's do more tests:

| Matrix         | Tolerance | Time (s) |
|----------------|-----------|----------|
| p2p-Gnutella30 | 1e-4      | 4.15     |
| p2p-Gnutella30 | 1e-25     | 4.06     |
| p2p-Gnutella30 | 1e-30     | 4.02     |

So, by lowering a bit the tolerance this method gets a little bit faster. That can be explain because it's using the sparse matrix properties. The other method, the one where we store the matrices, what really takes time is the creation of the  $z_j$  vector, and this is only created once. Once created, the method ends quickly.

You can see a time comparison in the figure 1.1.

### 1.4 Convergence

Computing the eigenvalues of A I get  $|\lambda_1| = 0.66$  and  $|\lambda_2| = 0.49$ . This means that their division ( $|\lambda_2|/|\lambda_1|$ ) is equal to 0.73, which is smaller than 1-0.15. In conclusion, it converges but it is not as fast as it could be, because the method converges slowly if there is an eigenvalue close in magnitude to the dominant eigenvalue.

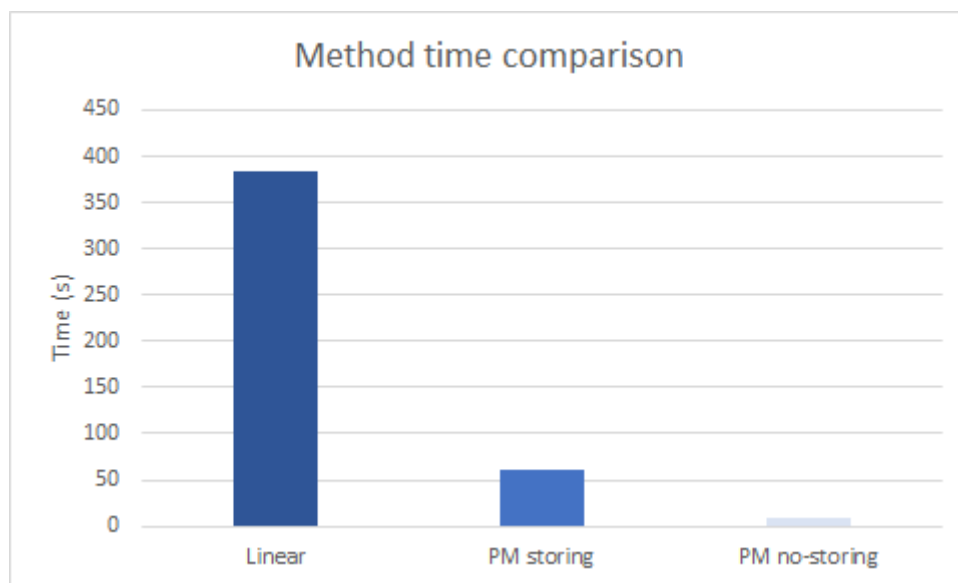


Figure 1.1: Comparison of the execution time of all 3 methods with a tolerance of  $1e-5$  and damping factor of 0.15